⇒ Indexing

DB stores data on a disk.

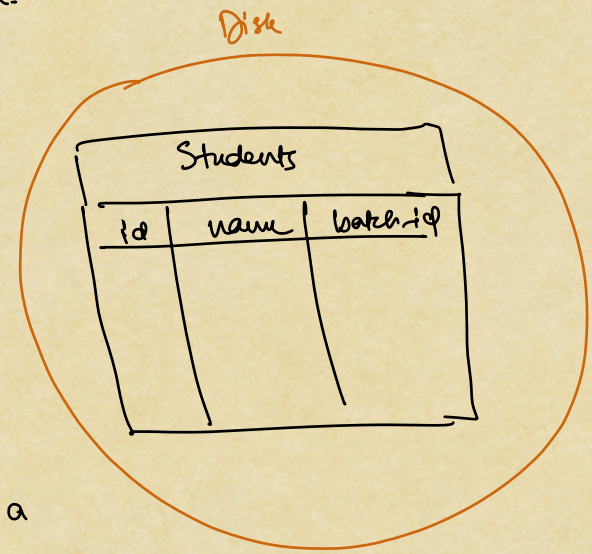

Select * from Students
where batch-id = 2;

* OS can directly operate on a disk
* Content from the disk is brought into memory and then appn works on that data.

⇒ If CPU were to directly work on the disk it will be wasting a lot of its time.

→ CPU makes an I/O call to disk for bringing the data to memory

→ while the data is being copied, it will do something else.

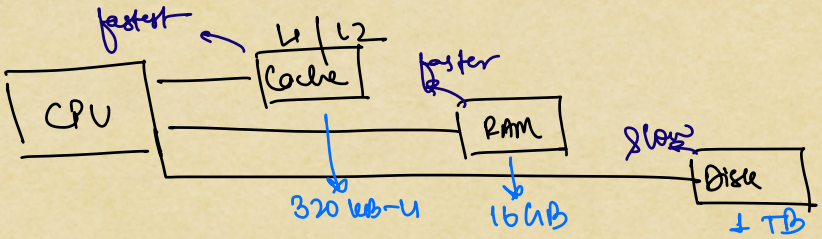

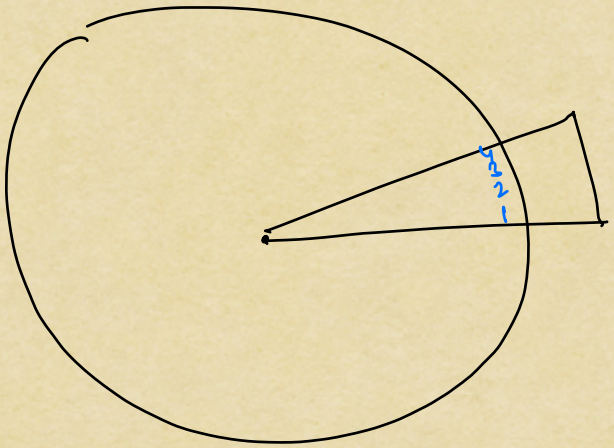

⇒ 1, 2, 3, 4

Query ⇒ 3

* A table is always sorted by its pk ⇒ default

Students → 1M rows

id	name	batch_id	psp	address
1	A	3	80	—
2	B	3	90	—
3	C	2	60	—
4	D	4	75	—
5	E	1	82	—

RAM & compare

Ram 2 compare

Ram 2 comp

RAM

RAM

Query ⇒ Select * from Students where batch_id = 3;

result ⇒ 20 Students

Total disk access = 1M

useless " " = 1M - 20

Query => Select * from Students where student_id = 3;

PK

Students → 1M rows

id	name	batch_id	psp	address
1	A	3	80	
2	B	3	90	
3	C	2	60	
4	D	4	75	
5	E	1	82	

memory addrs

#1 #2 #3 #4 #5

* because the table is sorted by pk, once we hit the condition (we get matching data) we stop & return
* still worst case can be 1M

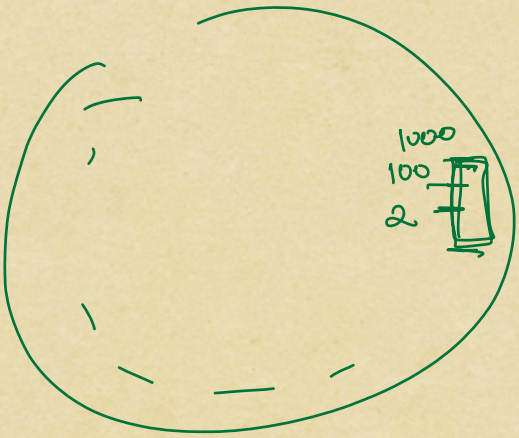

=> Create a table that stores id of a row and memory address of a row, sorted by id.

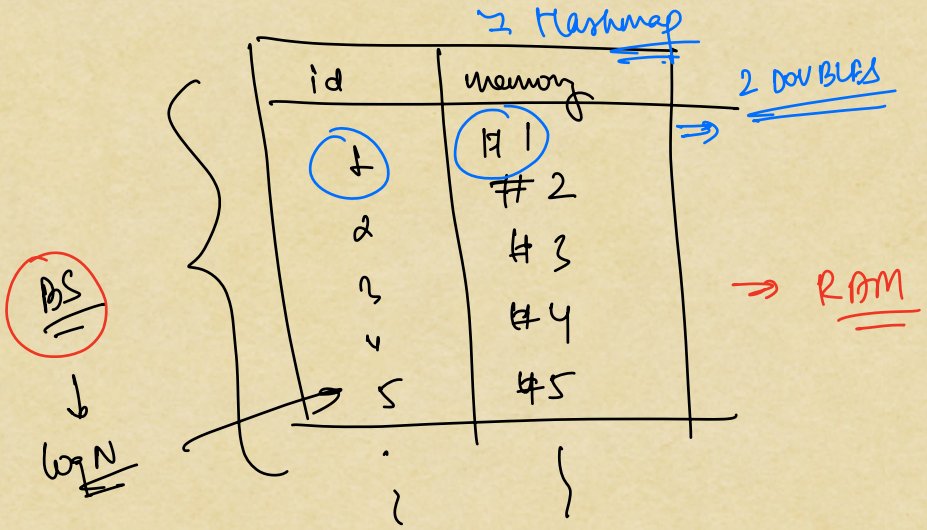

1M rows ⇒ Select * — — Where id = 15;

↓ pk

disk fetches

CASE 1 ⇒ No sorting | No address table ⇒ 1M

CASE 2 ⇒ Sorting | No address table ⇒ 15

CASE 3 ⇒ Sorting | address table ⇒ 1

* We want to reduce the no. of disk fetches.

eg ⇒

Students

id	name	batch_id	psp	phone No.
1	A	3	80	123 → #1
2	B	3	90	124 → #2
3	C	1	60	678 → #3
4	D	2	75	679 → #4
5	E	1	85	576 → #5

Select * from Student where phoneno = K;

↓

table → sorted by phone No.

phoNo.	addres
123	→ #1
124	→ #2
678	→ #3
679	→ #4
576	→ #5

⇒ Query ⇒ Select * from Student where batch_id = 3;

batch Id	address
3	#1
3	#2
1	#3
2	#4
1	#5

↓

sorted by batch-id

batch id	address
1	#3
1	#5
2	#4
3	#1
3	#2

ex ⇒ Select * from Students where psp is between 10 and 35;

psp	address
0	#1 #6 #7
10	10 9 8
20	2 4 5
30	31 63 32
40	52 62 64 47
50	
60	
70	
80	
90	
100	

→ ≥10 & <20

→ ≥20 & <30

→ ≥30 & <40

⇒ Indexes :

→ Prevent unnecessary disk fetches.

→ Leads to faster queries

→ Should we create indexes always?

↓

No. it depends.

⇒ C | U | D on table ⇒ indexes will also need to be updated.

⇒ if we do any write operations, index needs to update

⇒ indexing makes read faster & writes slower.

⇒ index tables are actually in disk, leads to increase in memory usage.

DS used for indexes ⇒ | B Tree | B+ Tree |

H/W ⇒ read about B Trees

⇒ When to create indexes:-

i) Do not create indexes at the beginning

ii) Do performance testing

iii) Create indexes for queries that are used a lot and are slow

iv) Create indexes by access patterns and not predictions.

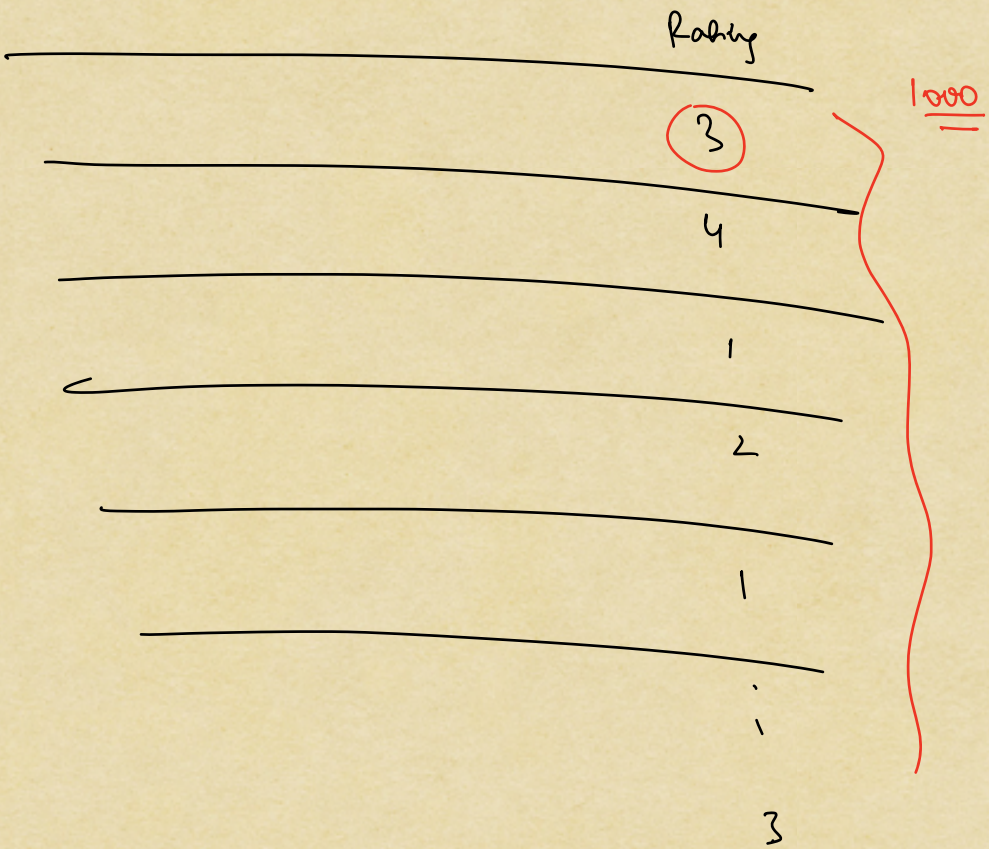

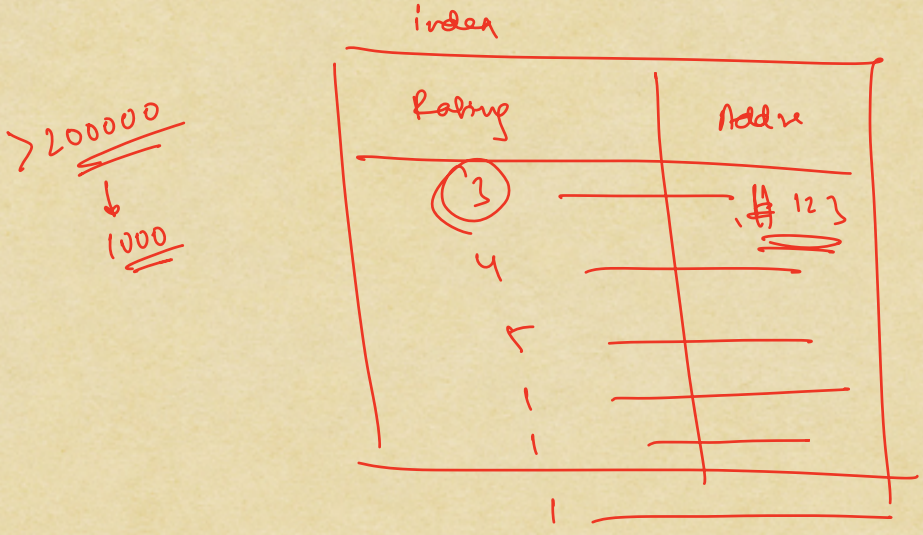
index
Rating
Addre
3
123
4
1
>200000
1000