

```
select *
from student
where name = 'Honny';
```

```
-- Go through the query stats.
-- Execution Plan
-- Full Table Scan
-- Optimizer calculates the Query Cost and based on the query cost, the decision will be taken
on which query to execute.
-- query cost = 3.75
```

```
select *
from student
where id = 10;
-- query cost = 1.0
```

```
-- A table is by default sorted by its PK.
-- Every table will have an index on the PK by default.
-- Single Row.
```

```
ALTER table student add index student_name_idx (name);
```

```
select *
from student
where name = 'Honny';
```

```
-- Query cost = 0.35
-- scan -> Non Unique Key lookup.
```

```
select o.order_id, o.shipped_date, c.customer_id, c.first_name, oi.product_id, oi.quantity,
       p.name, p.unit_price
from orders o
join customers c
on o.customer_id = c.customer_id
join order_items oi
on o.order_id = oi.order_id
join products p
on oi.product_id = p.product_id;
```

```
-- orders customers order_items products ->
```

```
-- Joins are very costly.
-- JOINS -> read
--
```

-- $O(N^3)$ -> Most optimal solution.

-- $O(N^2)$ -> $O(N \log N)$ -> $O(N)$

-- As a assignment, try to change the order of joins and then analyse the cost.

```
select o.order_id, o.shipped_date, c.customer_id, c.first_name, oi.product_id, oi.quantity
from orders o
join customers c
on o.customer_id = c.customer_id
join order_items oi
on o.order_id = oi.order_id;
```

```
select * from orders;
```

-- A join B

-- for every row A : $O(N)$

-- for every row of B : -> $O(N)$ worst

-- check the condition ON a.batch_id = b.batch_id

-- if the condition satisfies -> add in the output.