Agenda

1. Vertical v/s Horizontal Scaling
2. Load Balancers
   - Journey
3. Intro to Sharding
   - Hashing
   - Range Based
   - Core Hashing

Today

Next Class

If not attended prev class

" System Design and CN 101 "

- leave this class
- watch recording.

1. Patience

Recap of last Class

Case Study of Del.icio.us (2003)

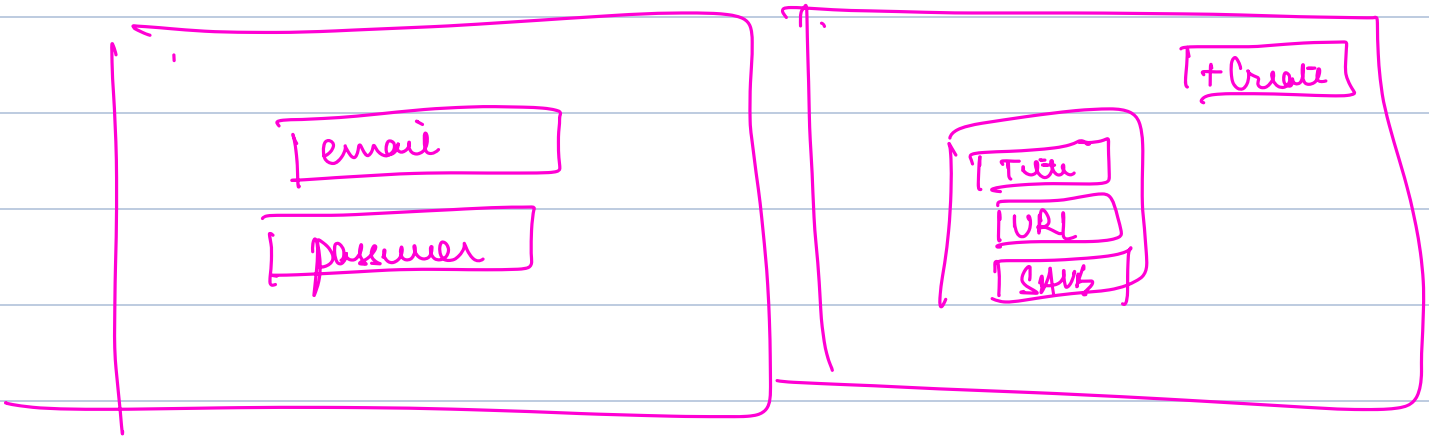

⇒ College Student

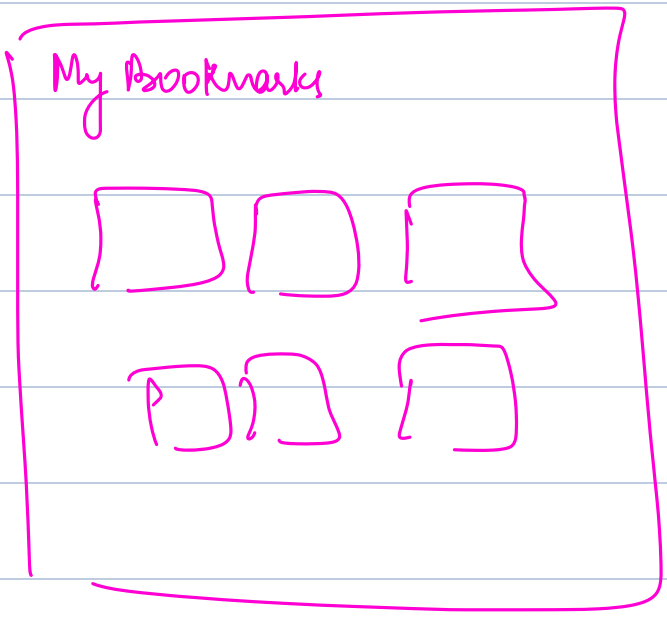

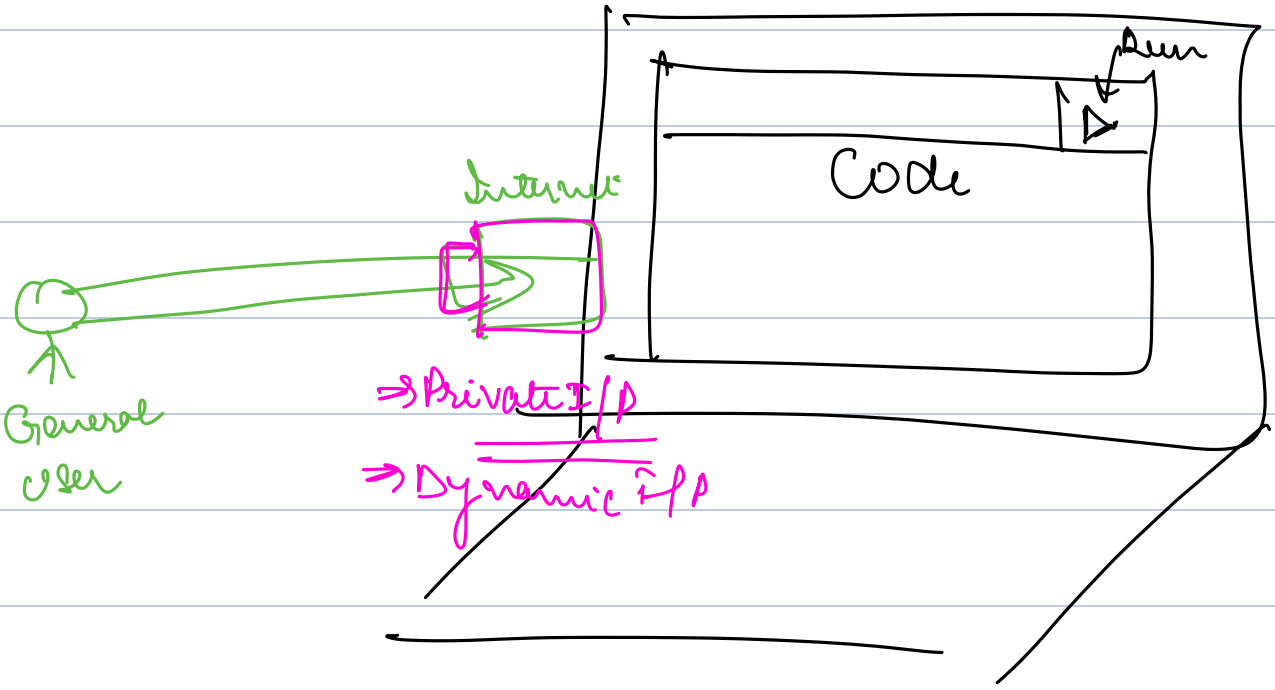

⇒ Private Dynamic I/P

Public Static I/P (to host website)

Pay Extra I/P

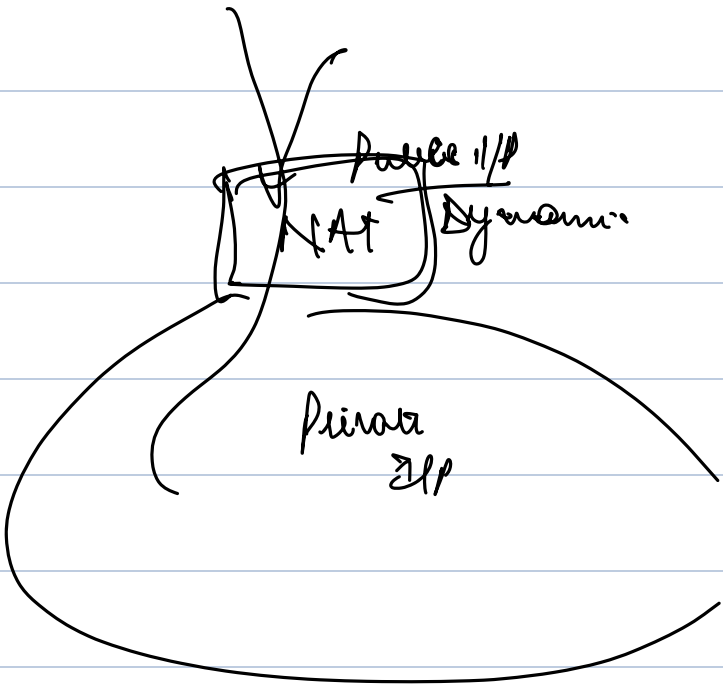

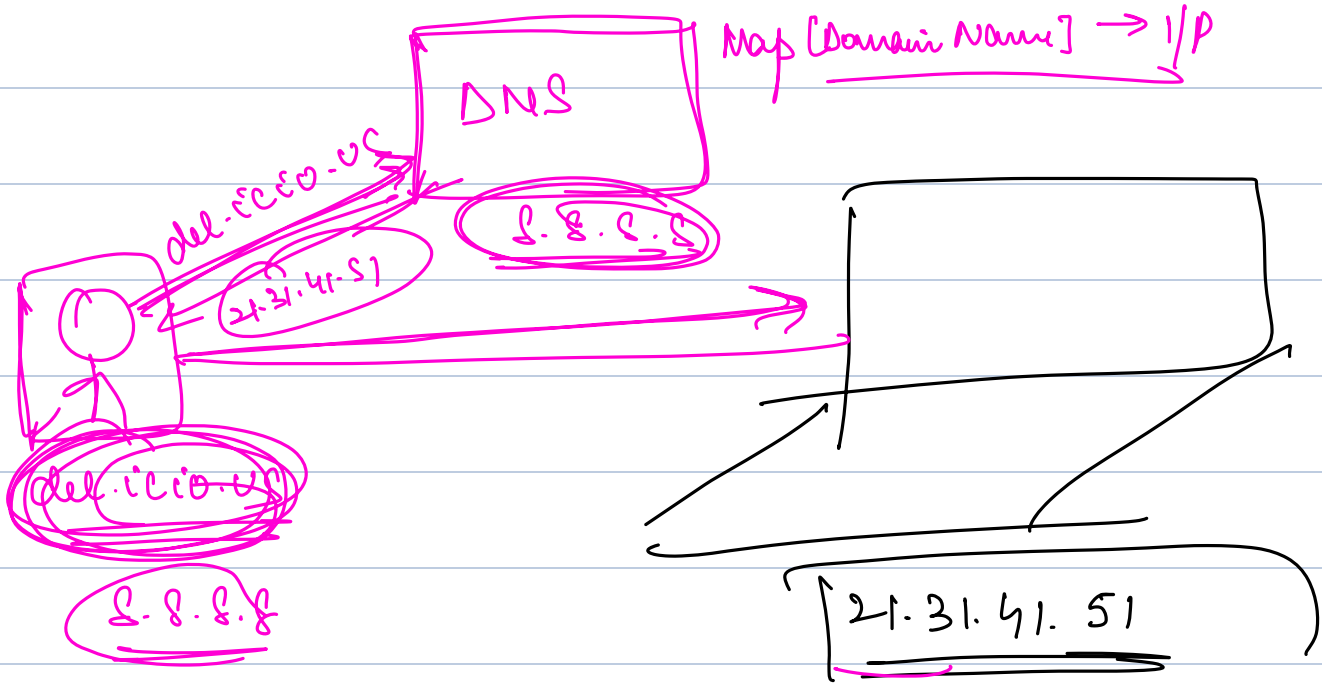

People like to type Names

21.31.41.51

Domain Name : google.com
fb.com

1. Bought a domain name (godaddy.com)
2. Got a public static IP connection
3. Wrote code
4. DNS updated the i/p of del.icio.us to the public static i/p

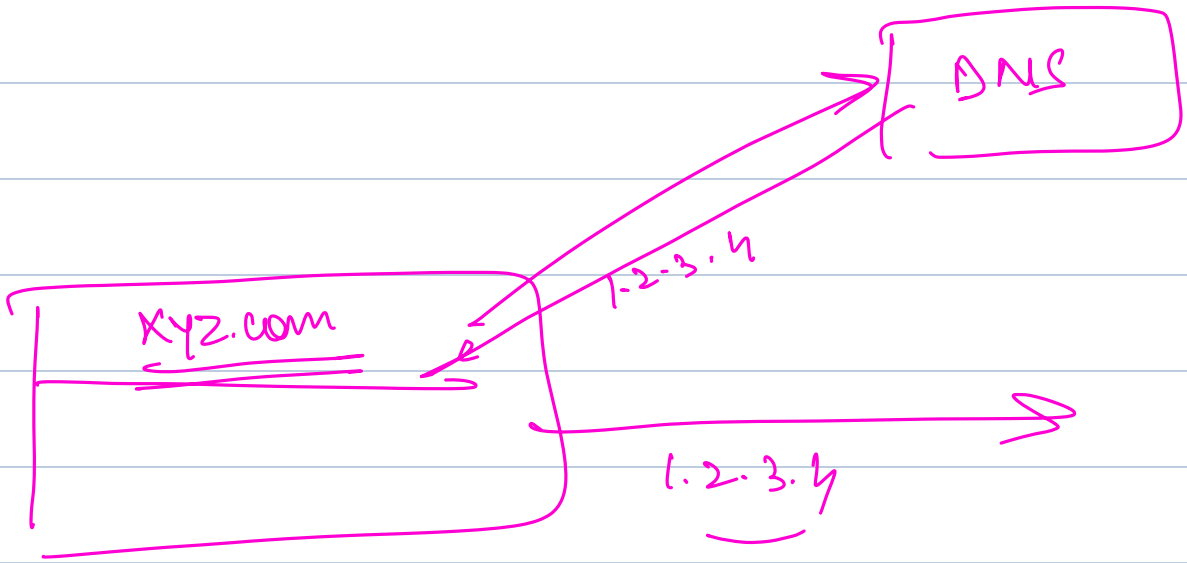

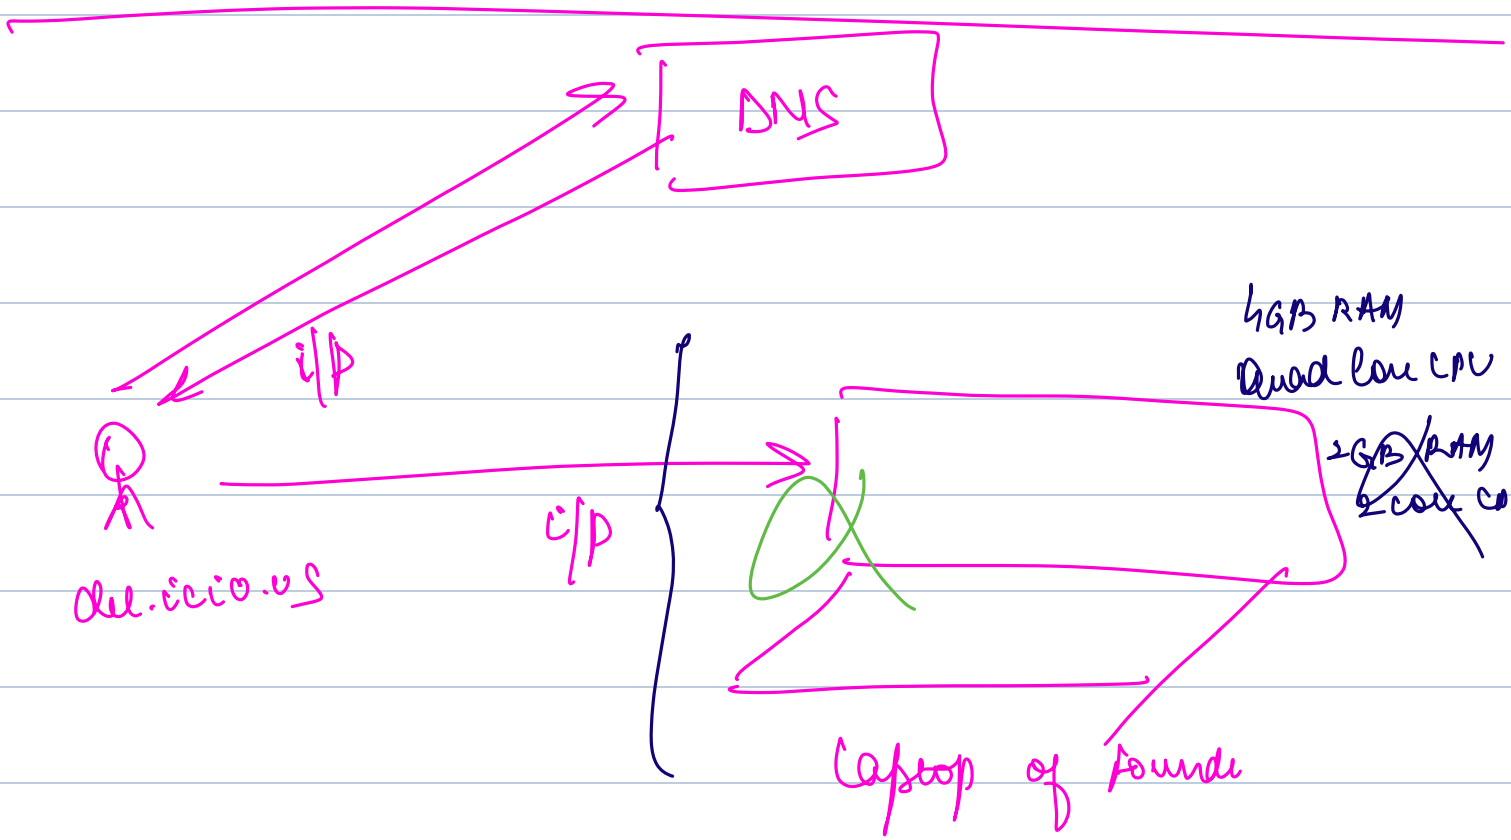

⇒ As website becomes more popular the laptop won't be able to handle all load / might crash.

Scaling: Ability of a website to handle more users

2 ways

Vertical Scaling

Still same amount of hardware But better hardware

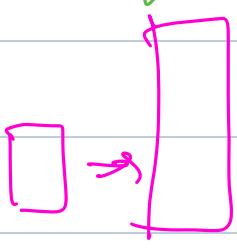

Pro

1.) Simple

Cons

1.) SPOF

2.) Not infinitely

Scalable

(Can't scale above a threshold)

3) Cost increases exponentially

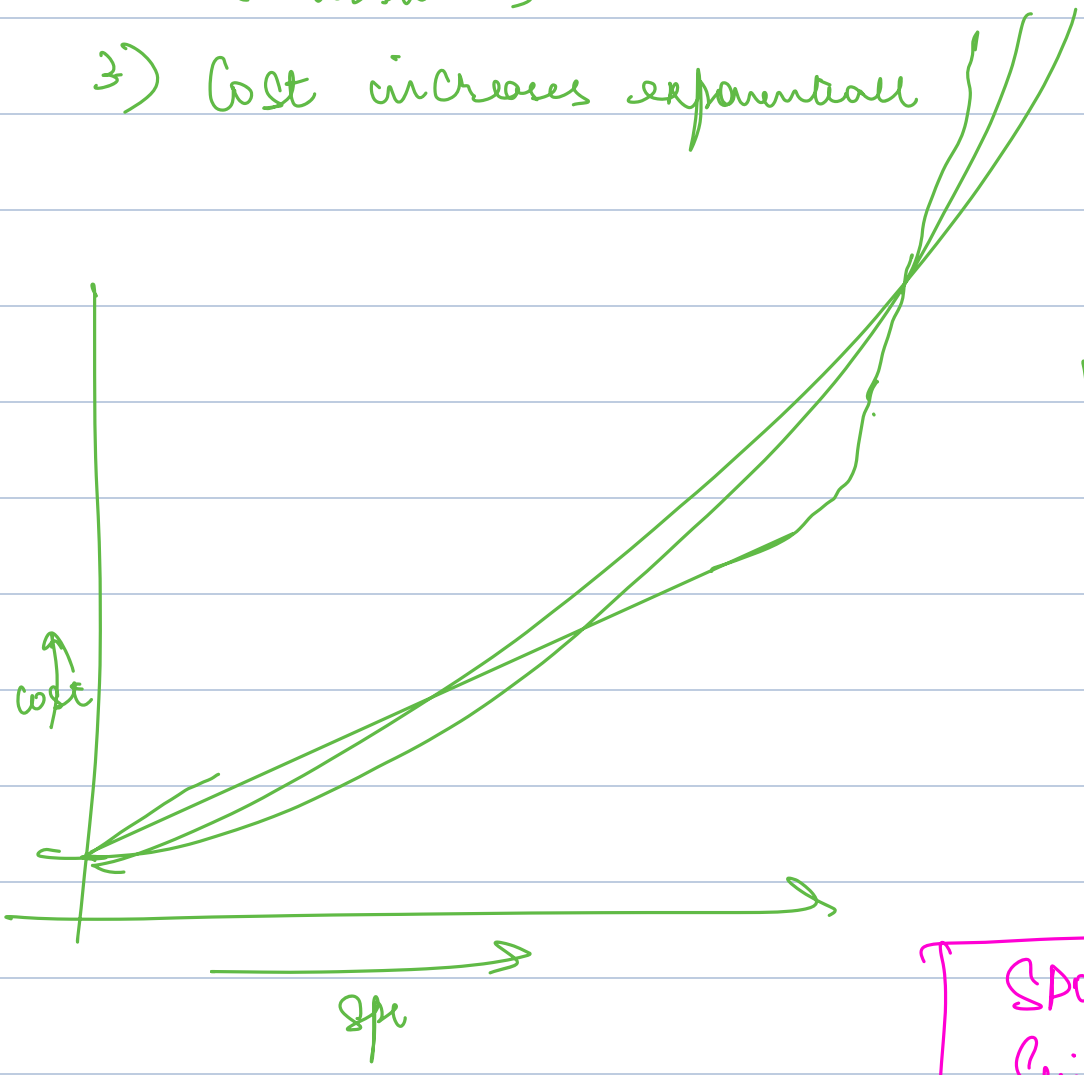

64GB ₹

128GB - 5/

SPOF
Single Point of Failure

② Horizontal Scaling

⇒ By getting more laptops, instead of only single

Pro:

→ Infinitely Scalable.

→ No SPOF

→ Linear cost increase

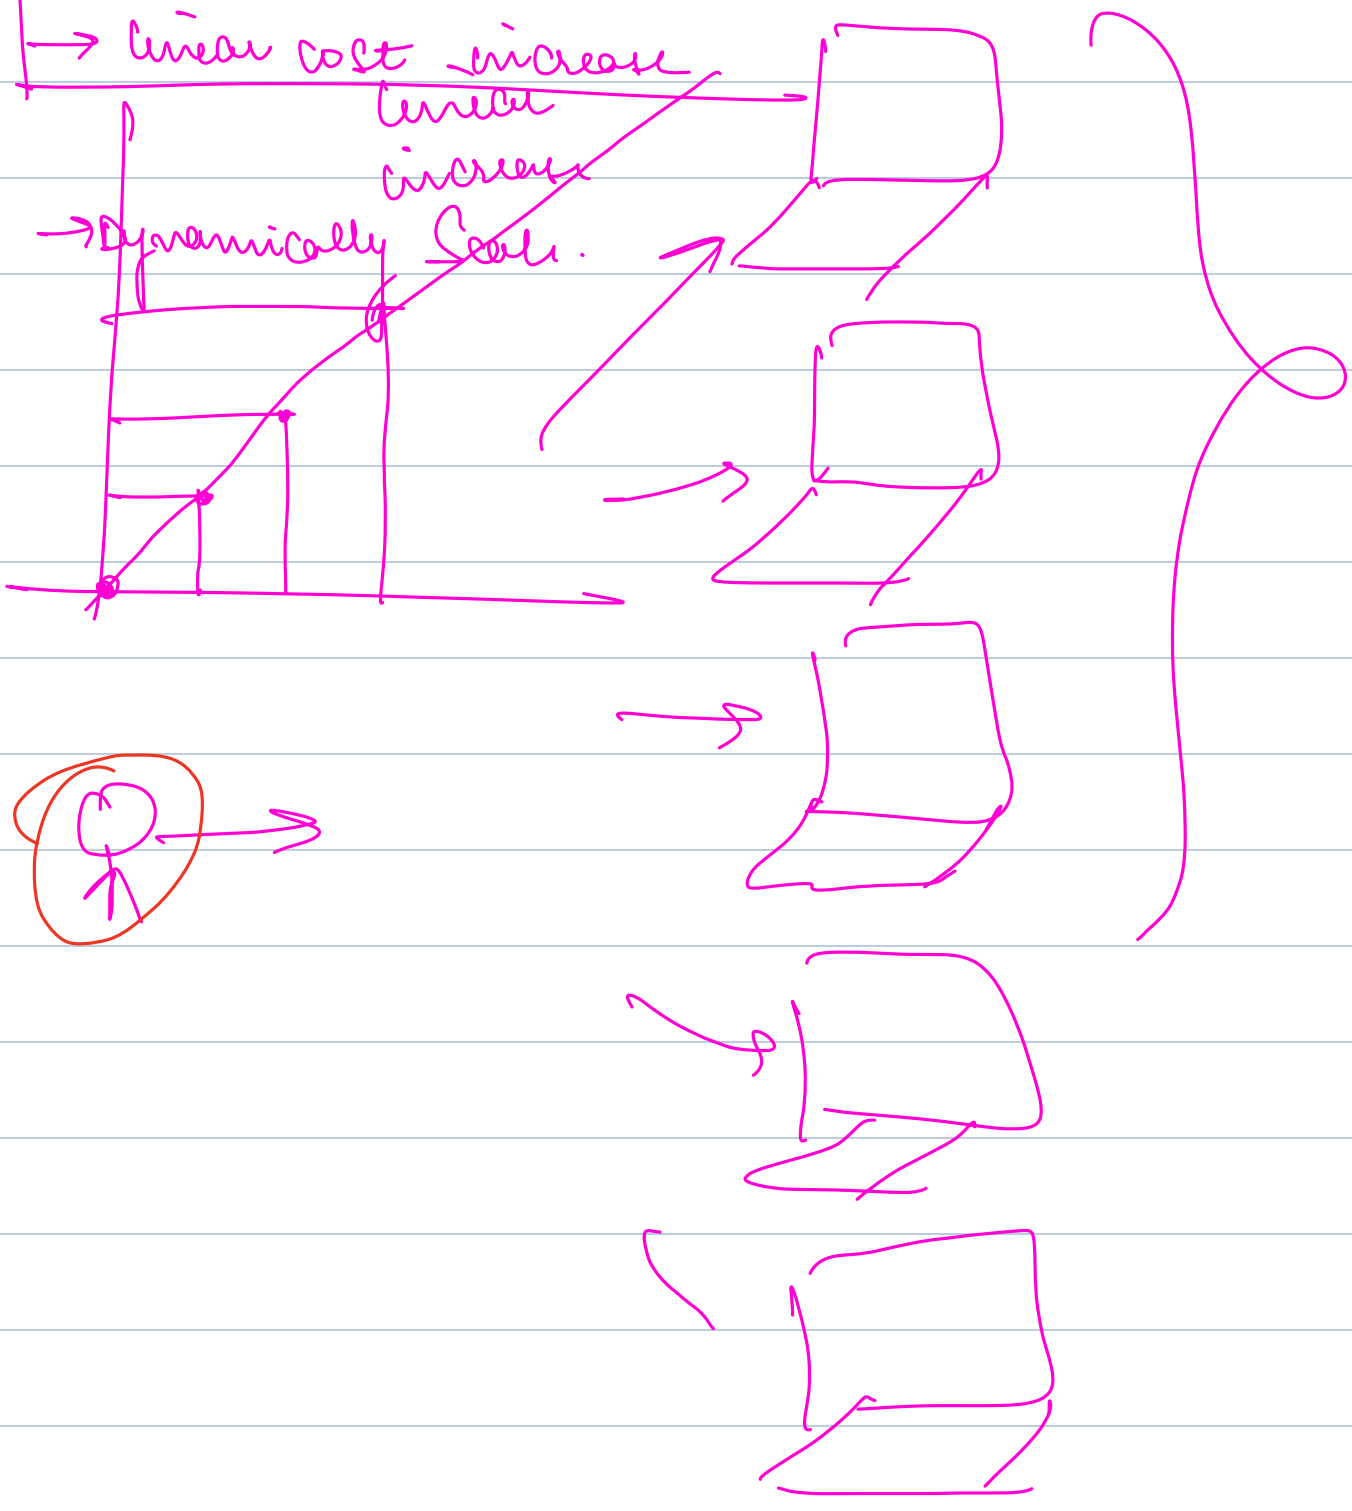

→ Dynamically Scale.

Cons

1.) Complexity of management
   - Maintain multiple machines
   - How to decide where to forward a req.

KISS

Keep it simple, Stupid

How does horizontal Scaling work in reality

Load Balancer

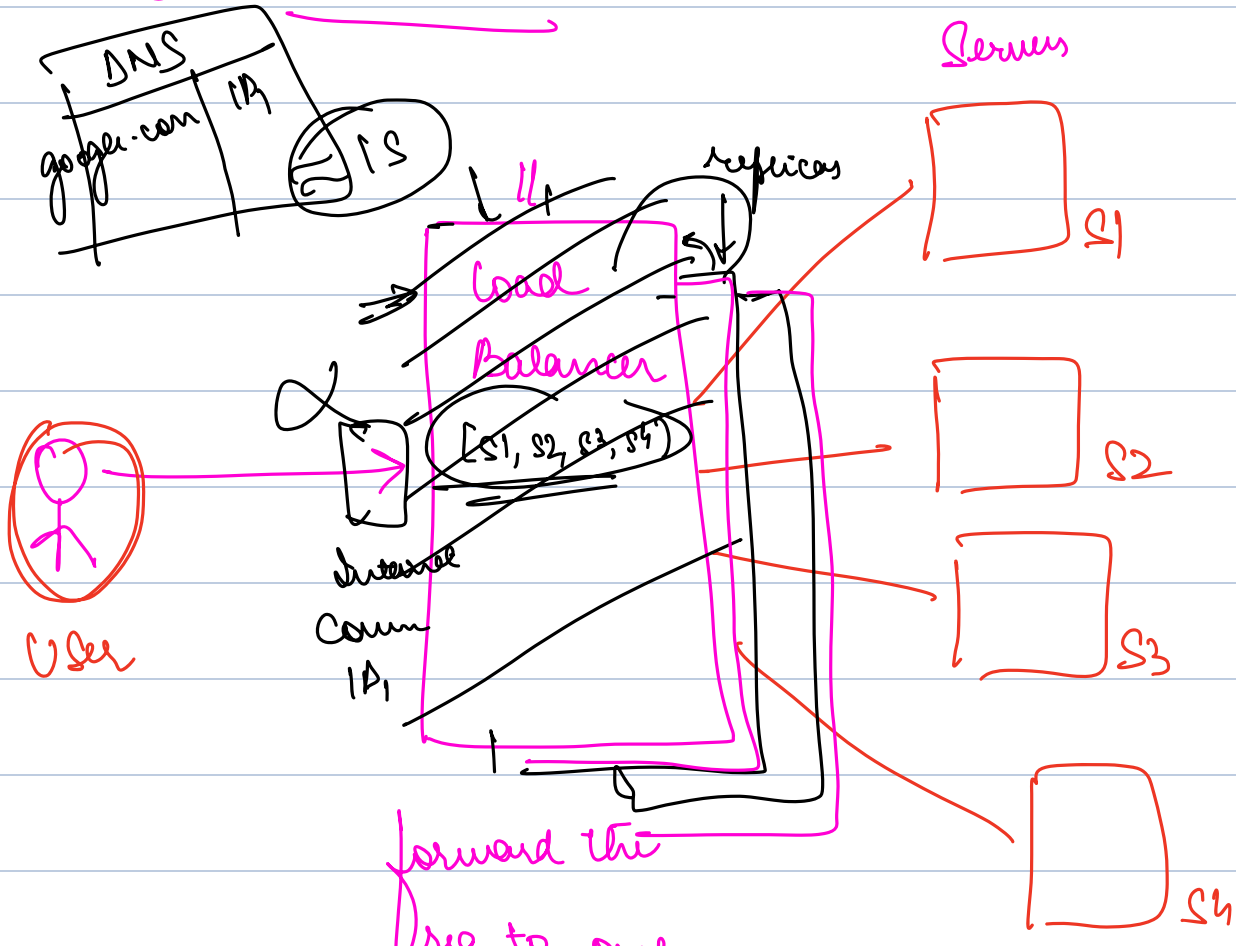

LB: distribute the req amongst servers in a uniform way

(So each server is almost equally loaded

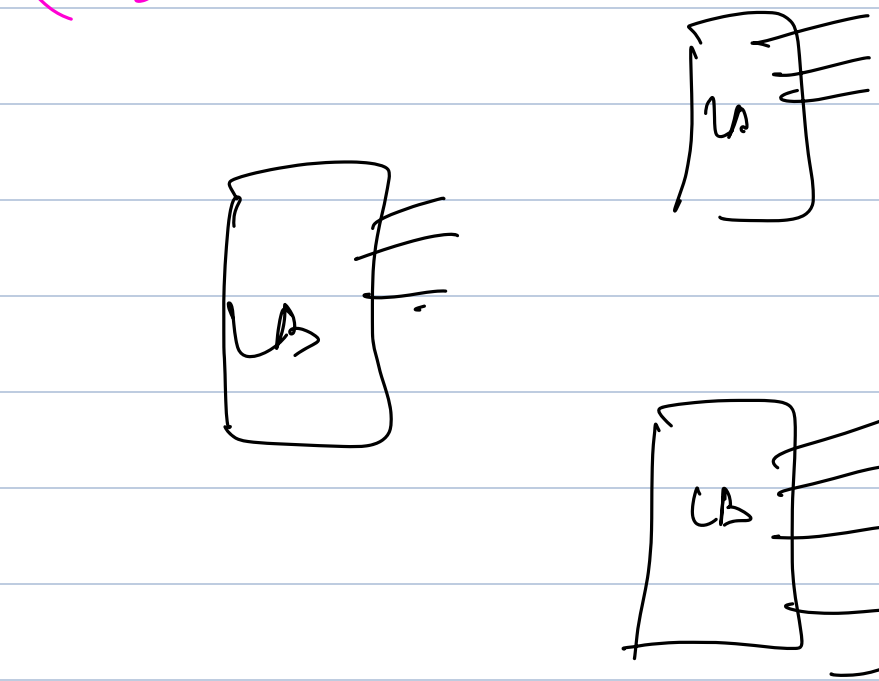

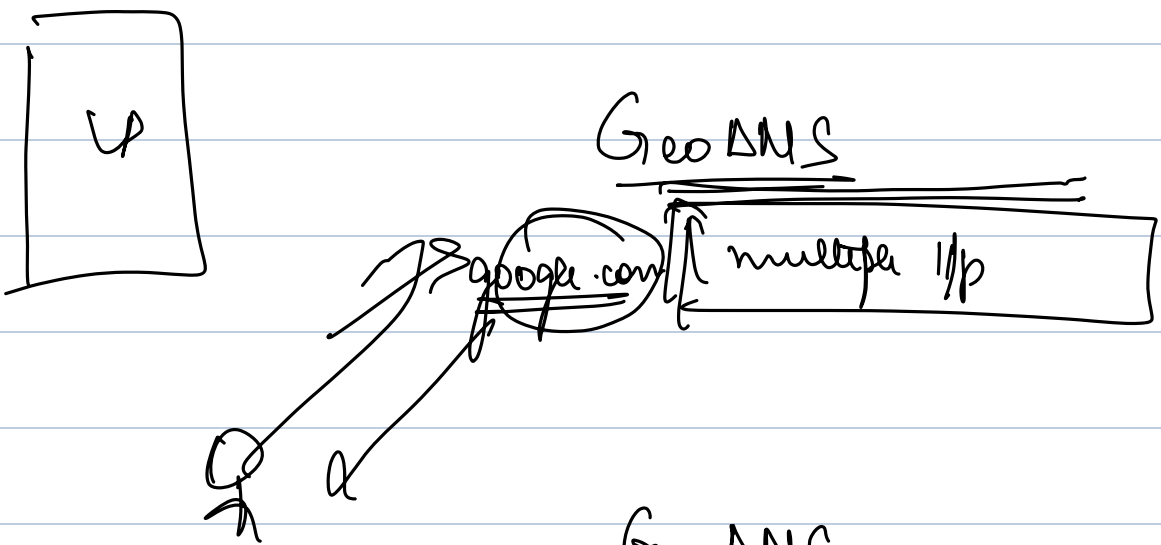

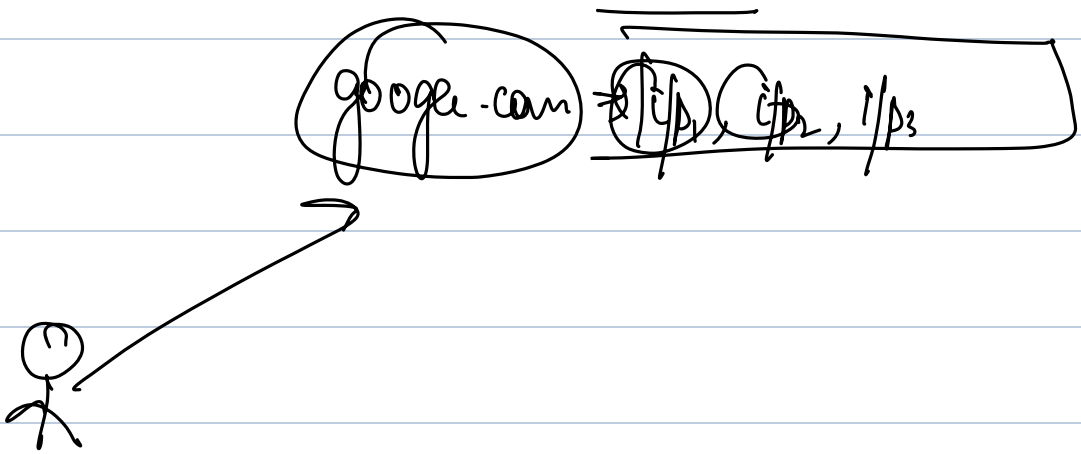

How will LB work

- How will LB get to know servers
- How will LB know of server dead
- How will LB decide where to forward req.

How will LB get to know a server

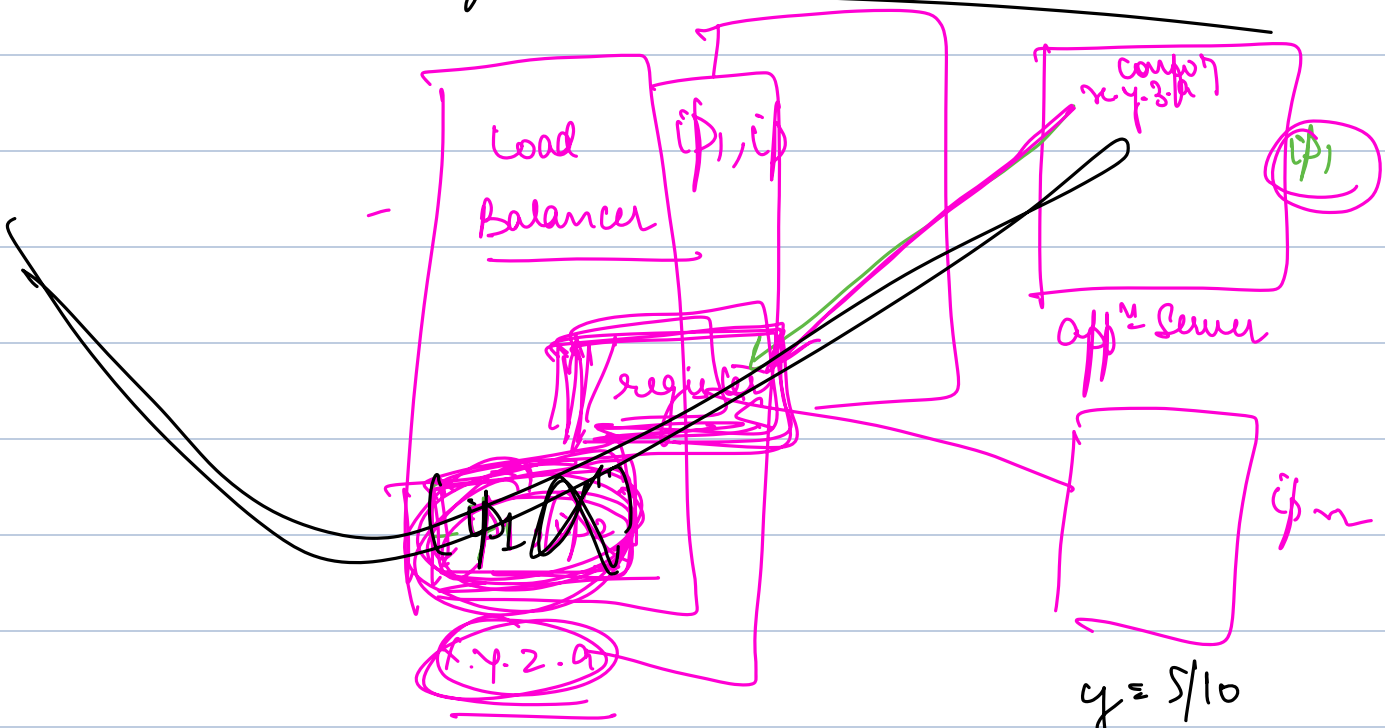

x.y.2.a/register

How will LB know that app^n server has died

1.) HealthCheck
2.) Heart Beat

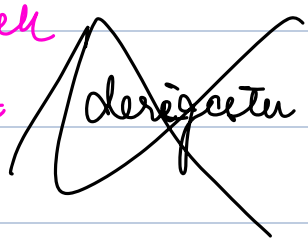

Normal Days

Parents care of your health
they calls you

Health Check: every "x" sec LB will send a req to every server asking if it's okay

(i) if Not okay → delegate
(ii) if server doesn't reply for 3 times
→ Assume it has died
→ remove it

/health

Heart Beat :

every server has to send an "OK" to LB every 'X' sec (10 sec)

if LB doesn't get consecutive 3 OK, it will assume dead

How will LB forward Req

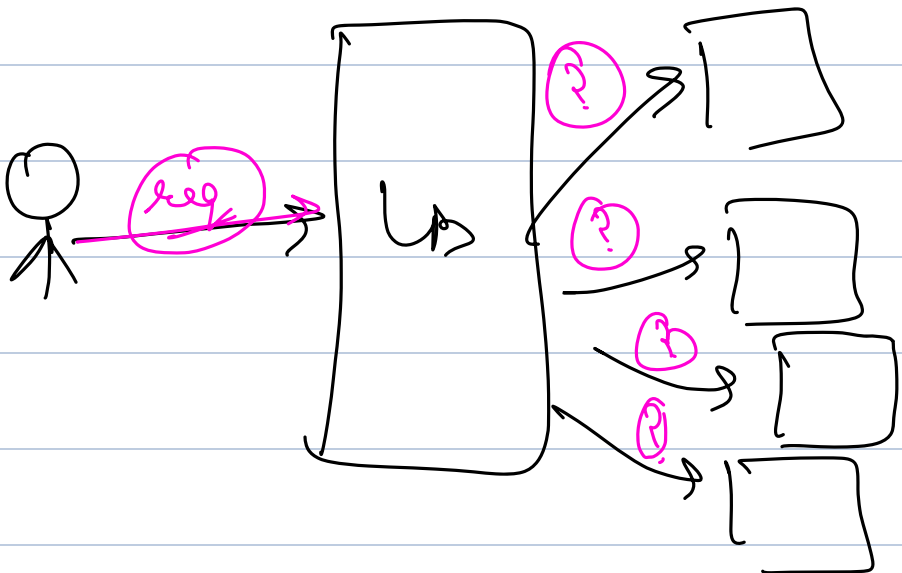

Virus from 3 IDIOTS "as expected"

Approach 1

LB asks each server about their current CPU/RAM usage forwards to least one.

→ Terrible.

→ Make the req slow

Approach 2: Round Robin

(principle of probabilities)

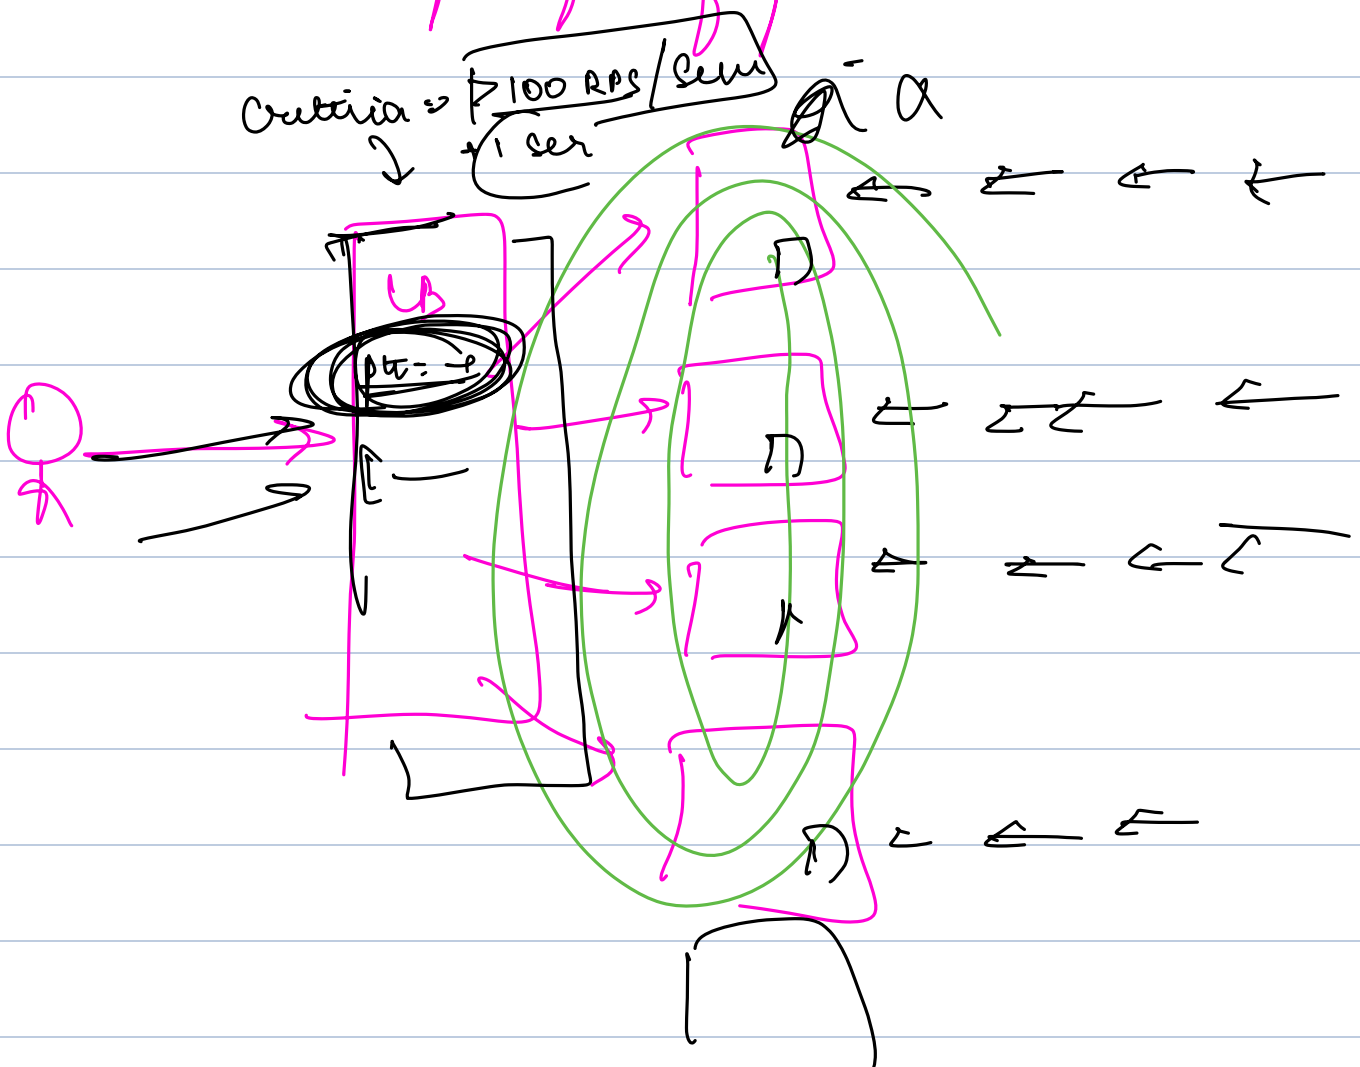

Built ⇒ Registering

Death ⇒ Health Check

HeartBeat

Forward ⇒ RR

- RPS (Requests Per Sec)
- Av. GPU usage
- Av RAM usage

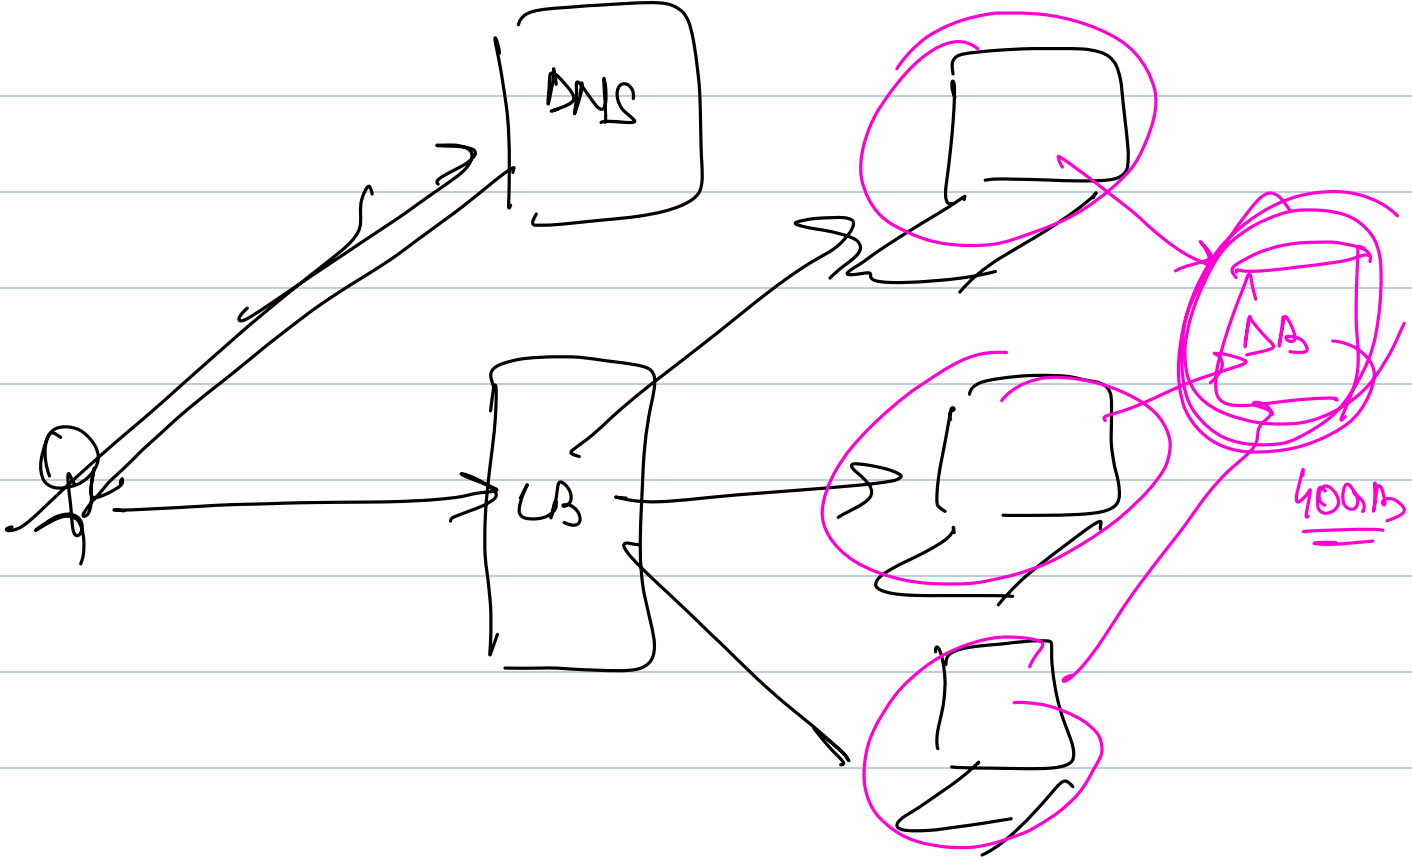

2003

HDD ⇒ 40 GB

2004

72GB HDD

400M users

every user ≈ 20 BM

⇒ Total BM: $400 \times 20 \times 10^6$

⇒ 8×10^9 ⇒ 8B BM

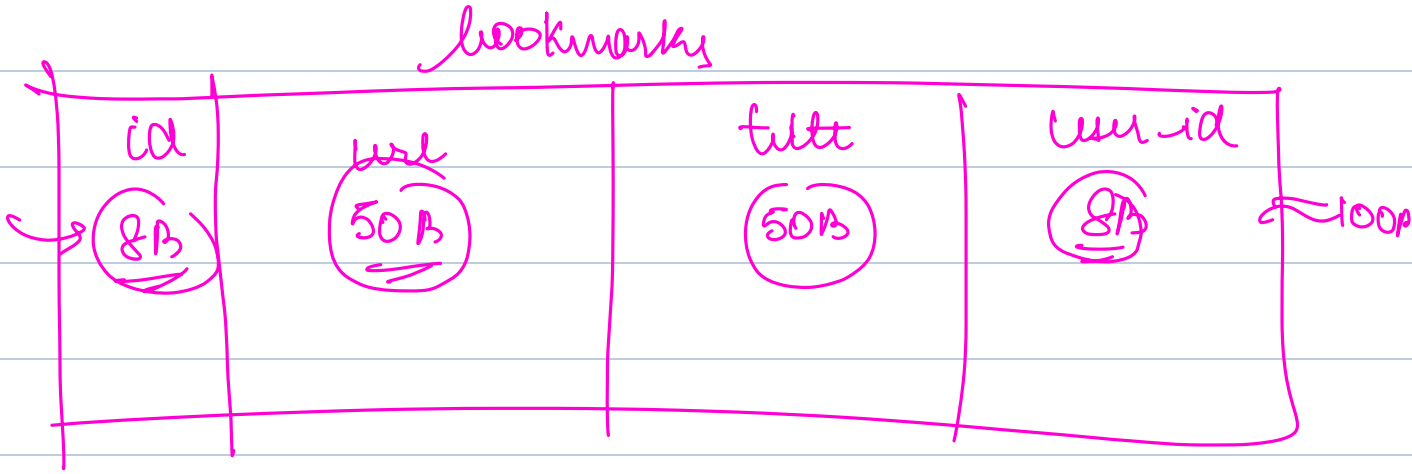

⇒ Total Size

→ $8 \times 10^9 \times 100$ Bytes

→ 800×10^9 Bytes

→ 800 GB

⇒ Unfortunately 1 machine can't have all data

→ Split Data

→ Sharding